

# Constraints and Triggers

# Outline

1. Introduction
2. Constraints
3. Triggers

# Constraints and Triggers

- A *constraint* is a relationship among data elements that the DBMS is required to enforce.
  - *Example*: key constraints.
- *Triggers* are only executed when a specified condition occurs, e.g., insertion of a tuple.
  - Easier to implement than complex constraints.

# Outline

1. Introduction
- 2. Constraints**
3. Triggers

# Kinds of Constraints

- **Keys** (primary key or unique key).
- **Foreign-key**, or referential-integrity.
- **Value-based** constraints.
  - Constrain values of a particular attribute.
- **Tuple-based** constraints.
  - Relationship among components.
- **Assertions**: any SQL boolean expression.

# Review: Keys

- An attribute or list of attributes may be declared PRIMARY KEY or UNIQUE.
- Either says that no two tuples of the relation may agree in all the attribute(s) on the list.
- There are a few distinctions to be mentioned later.

# PRIMARY KEY vs. UNIQUE

1. There can be only one PRIMARY KEY for a relation, but several UNIQUE attributes.
2. No attribute of a PRIMARY KEY can ever be NULL in any tuple. But attributes declared UNIQUE may have NULL's, and there may be several tuples with NULL.

# Review: Declaring a Relation

- Simplest form is:

```
CREATE TABLE <name> (  
    <list of elements>  
);
```

- To delete a relation:

```
DROP TABLE <name>;
```



# Review: Create Table

```
CREATE TABLE Sells (  
    bar        CHAR(20) ,  
    beer       VARCHAR(20) ,  
    price      REAL  
);
```

# Declaring Single-Attribute Keys

- Place PRIMARY KEY or UNIQUE after the type in the declaration of the attribute.
- Example:

```
CREATE TABLE Beers (  
    name        CHAR(20)  UNIQUE,  
    manf        CHAR(20)  
);
```

# Declaring Multiattribute Keys

- A key declaration can also be another element in the list of elements of a `CREATE TABLE` statement.
- This form is essential if the key consists of more than one attribute.
  - May be used even for one-attribute keys.

# Example: Multiattribute Key

- The bar and beer together are the key for Sells:

```
CREATE TABLE Sells (  
    bar          CHAR(20) ,  
    beer         VARCHAR(20) ,  
    price        REAL ,  
    PRIMARY KEY (bar, beer)  
);
```

# Foreign Keys

- Values appearing in attributes of one relation must appear together in certain attributes of another relation.
- **Example:** in **Sells(bar, beer, price)**, we might expect that a beer value also appears in Beers.name .

# Expressing Foreign Keys

- Use keyword REFERENCES, either:
  1. After an attribute (for one-attribute keys).
  2. As an element of the schema:  
FOREIGN KEY (<list of attributes>)  
REFERENCES <relation> (<attributes>)
- Referenced attributes must be declared PRIMARY KEY or UNIQUE.

# Example: With Attribute

```
CREATE TABLE Beers (  
    name      CHAR(20) PRIMARY KEY,  
    manf      CHAR(20) );  
  
CREATE TABLE Sells (  
    bar       CHAR(20) ,  
    beer      CHAR(20) REFERENCES Beers(name) ,  
    price     REAL );
```

# Example: As Schema Element

```
CREATE TABLE Beers (  
    name      CHAR(20) PRIMARY KEY,  
    manf      CHAR(20) );  
  
CREATE TABLE Sells (  
    bar       CHAR(20),  
    beer      CHAR(20),  
    price     REAL,  
    FOREIGN KEY (beer) REFERENCES  
        Beers (name) );
```



# Enforcing Foreign-Key Constraints

- If there is a foreign-key constraint from relation  $R$  to relation  $S$ , two violations are possible:
  1. An insert or update to  $R$  introduces values not found in  $S$ .
  2. A deletion or update to  $S$  causes some tuples of  $R$  to “dangle.”

# Actions Taken --- (1)

- **Example:** suppose  $R = \text{Sells}$ ,  $S = \text{Beers}$ .
- An insert or update to **Sells** that introduces a nonexistent beer must be rejected.
- A deletion or update to **Beers** that removes a beer value found in some tuples of **Sells** can be handled in three ways (next slide).

# Actions Taken --- (2)

1. *Default* : Reject the modification.
2. *Cascade* : Make the same changes in Sells.
  - Deleted beer: delete Sells tuple.
  - Updated beer: change value in Sells.
3. *Set NULL* : Change the beer to NULL.

# Example: Cascade

- Delete the Bud tuple from Beers:
  - Then delete all tuples from Sells that have beer = 'Bud'.
- Update the Bud tuple by changing 'Bud' to 'Budweiser':
  - Then change all Sells tuples with beer = 'Bud' to beer = 'Budweiser'.

## Example: Set NULL

- Delete the Bud tuple from Beers:
  - Change all tuples of Sells that have beer = 'Bud' to have beer = NULL.
- Update the Bud tuple by changing 'Bud' to 'Budweiser':
  - Same change as for deletion.

# Choosing a Policy

- When we declare a foreign key, we may choose policies SET NULL or CASCADE independently for deletions and updates.
- Follow the foreign-key declaration by:  
ON [UPDATE, DELETE][SET NULL|CASCADE]
- Two such clauses may be used.
- Otherwise, the default (reject) is used.

# Example: Setting Policy

```
CREATE TABLE Sells (  
    bar    CHAR(20),  
    beer   CHAR(20),  
    price  REAL,  
    FOREIGN KEY (beer)  
        REFERENCES Beers (name)  
        ON DELETE SET NULL  
        ON UPDATE CASCADE  
);
```

# Attribute-Based Checks

- Constraints on the value of a particular attribute.
- Add CHECK(<condition>) to the declaration for the attribute.
- The condition may use the name of the attribute, but **any other relation or attribute name must be in a subquery.**



# Example: Attribute-Based Check

```
CREATE TABLE Sells (  
    bar    CHAR(20),  
    beer   CHAR(20)    CHECK ( beer IN  
        (SELECT name FROM Beers) ),  
    price  REAL CHECK ( price <= 5.00 )  
);
```

# Timing of Checks

- Attribute-based checks are performed only when a value for that attribute is inserted or updated.
  - **Example:** `CHECK (price <= 5.00)` checks every new price and rejects the modification (for that tuple) if the price is more than \$5.
  - **Example:** `CHECK (beer IN (SELECT name FROM Beers))` not checked if a beer is deleted from Beers (unlike foreign-keys).

# Tuple-Based Checks

- CHECK (<condition>) may be added as a relation-schema element.
- The condition may refer to any attribute of the relation.
  - But other attributes or relations require a subquery.
- Checked on insert or update only.

# Example: Tuple-Based Check

- Only Joe's Bar can sell beer for more than \$5:

```
CREATE TABLE Sells (  
    bar          CHAR(20) ,  
    beer         CHAR(20) ,  
    price        REAL,  
    CHECK (bar = 'Joe' 's Bar' OR  
           price <= 5.00)  
);
```

# Assertions

- These are database-schema elements, like relations or views.
- Defined by:

```
CREATE ASSERTION <name>  
CHECK (<condition>);
```

- Condition may refer to any relation or attribute in the database schema.

# Example: Assertion

- In `Sells(bar, beer, price)`, no bar may charge an average of more than \$5.


```
CREATE ASSERTION NoRipoffBars CHECK (  
  NOT EXISTS (  

```

```
    SELECT bar FROM Sells  
    GROUP BY bar  
    HAVING 5.00 < AVG(price)
```

```
  ));
```

Bars with an  
average price  
above \$5



# Example: Assertion

- In **Drinkers(name, addr, phone)** and **Bars(name, addr, license)**, there cannot be more bars than drinkers.

```
CREATE ASSERTION FewBar CHECK (  
    (SELECT COUNT(*) FROM Bars) <=  
    (SELECT COUNT(*) FROM Drinkers)  
);
```

# Timing of Assertion Checks

- In principle, we must check every assertion after every modification to any relation of the database.
- A clever system can observe that only certain changes could cause a given assertion to be violated.
  - **Example:** No change to Beers can affect FewBar. Neither can an insertion to Drinkers.



# Outline

1. Introduction
2. Constraints
- 3. Triggers**

# Triggers: Motivation

- Assertions are powerful, but the DBMS often can't tell when they need to be checked.
- Attribute- and tuple-based checks are checked at known times, but are not powerful.
- Triggers let the user decide when to check for any condition.

# Event-Condition-Action Rules

- Another name for “trigger” is *ECA rule*, or *event-condition-action* rule.
- *Event*: typically a type of database modification, e.g., “insert on Sells.”
- *Condition*: Any SQL boolean-valued expression.
- *Action*: Any SQL statements.

# Preliminary Example: A Trigger

- Instead of using a foreign-key constraint and rejecting insertions into `Sells(bar, beer, price)` with unknown beers, a trigger can add that beer to `Beers`, with a NULL manufacturer.

# Example: Trigger Definition

```
CREATE TRIGGER BeerTrig
  AFTER INSERT ON Sells
  REFERENCING NEW ROW AS NewTuple
  FOR EACH ROW
  WHEN (NewTuple.beer NOT IN
        (SELECT name FROM Beers))
  INSERT INTO Beers(name)
  VALUES(NewTuple.beer);
```

The event

The condition

The action

# Options: CREATE TRIGGER

- CREATE TRIGGER <name>
- Or:

CREATE OR REPLACE TRIGGER <name>

- Useful if there is a trigger with that name and you want to modify the trigger.

# Options: The Event

- AFTER can be BEFORE.
  - Also, INSTEAD OF, if the relation is a view.
    - A clever way to execute view modifications: have triggers translate them to appropriate modifications on the base tables.
- INSERT can be DELETE or UPDATE.
  - And UPDATE can be UPDATE . . . ON a particular attribute.

# Options: FOR EACH ROW

- Triggers are either “row-level” or “statement-level.”
- FOR EACH ROW indicates row-level; its absence indicates statement-level.
- *Row level triggers* : execute once for each modified tuple.
- *Statement-level triggers* : execute once for a SQL statement, regardless of how many tuples are modified.



# Options: REFERENCING

- INSERT statements imply a new tuple (for row-level) or new table (for statement-level).
  - The “table” is the set of inserted tuples.
- DELETE implies an old tuple or table.
- UPDATE implies both.
- Refer to these by  
[NEW | OLD][TUPLE | TABLE] AS <name>

# Options: The Condition

- Any boolean-valued condition.
- Evaluated on the database as it would exist before or after the triggering event, depending on whether BEFORE or AFTER is used.
  - But always before the changes take effect.
- Access the new/old tuple/table through the names in the REFERENCING clause.

# Options: The Action

- There can be more than one SQL statement in the action.
  - Surround by BEGIN . . . END if there is more than one.
- But queries make no sense in an action, so we are really limited to modifications.

# Another Example

- Using `Sells(bar, beer, price)` and a unary relation `RipoffBars(bar)`, maintain a list of bars that raise the price of any beer by more than \$1.

# The Trigger

```
CREATE TRIGGER PriceTrig
```

```
AFTER UPDATE OF price ON Sells
```

The event –  
only changes  
to prices

```
REFERENCING
```

```
OLD ROW AS ooo
```

```
NEW ROW AS nnn
```

Updates let us  
talk about old  
and new tuples

We need to consider  
each price change

Condition:  
a raise in  
price > \$1

```
FOR EACH ROW
```

```
WHEN(nnn.price > ooo.price + 1.00)
```

```
INSERT INTO RipoffBars  
VALUES(nnn.bar);
```

When the price change  
is great enough, add  
the bar to RipoffBars