

# Operating Systems Memory Management

Christian Khoury

# Objectives

- Understand the various memory-management mechanisms
  - Contiguous
  - Paging
  - Segmentation

# Contents

- **Definitions**
- Contiguous Allocation
- Paging
- Segmentation
- Virtual Memory

# Definitions

- Logical/Virtual Address
  - Generated by the CPU
- Physical Address
  - Generated by the MMU

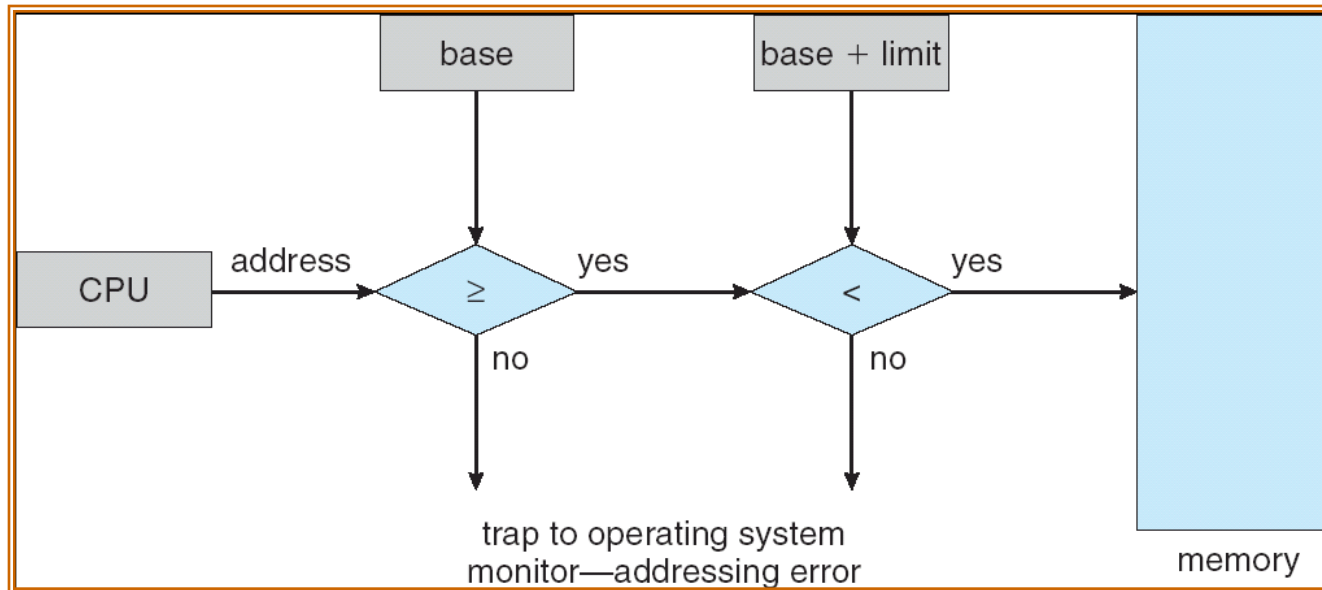
# Contents

- Definitions
- **Contiguous Allocation**
- Paging
- Segmentation
- Virtual Memory

# Contiguous Allocation (1/5)

- Base register contains value of smallest physical address
- Limit register contains range of logical addresses – each logical address must be less than the limit register
- MMU maps logical address *dynamically*

# Contiguous Allocation (2/5)



# Contiguous Allocation (3/5)

- Multiple-partition allocation
  - Hole – block of available memory; holes of various size are scattered throughout memory
  - When a process arrives, it is allocated memory from a hole large enough to accommodate it
  - Operating system maintains information about:
    - a) allocated partitions    b) free partitions (hole)



# Contiguous Allocation (4/5)

- Dynamic memory allocation strategies
  - First Fit
    - First on route “big enough” hole is used
  - Best Fit
    - Allocate the smallest hole that is big enough
    - Smallest leftover holes
  - Worst Fit
    - Allocate the largest hole that is big enough
    - Largest leftover holes

# Contiguous Allocation (5/5)

- External Fragmentation
  - Processes are loaded and removed from memory
  - Free memory (holes) is broken into little pieces (Fragments)
  - May have enough global memory for a request but noncontiguous !
- Solutions
  - Compaction : Not always possible, expensive
  - Paging, Segmentation

# Contents

- Definitions
- Contiguous Allocation
- **Paging**
- Segmentation
- Virtual Memory

# Paging (1/11)

- Memory management scheme that permits the physical address space to be noncontiguous/fragmented
- Need to map logical addresses into physical ones

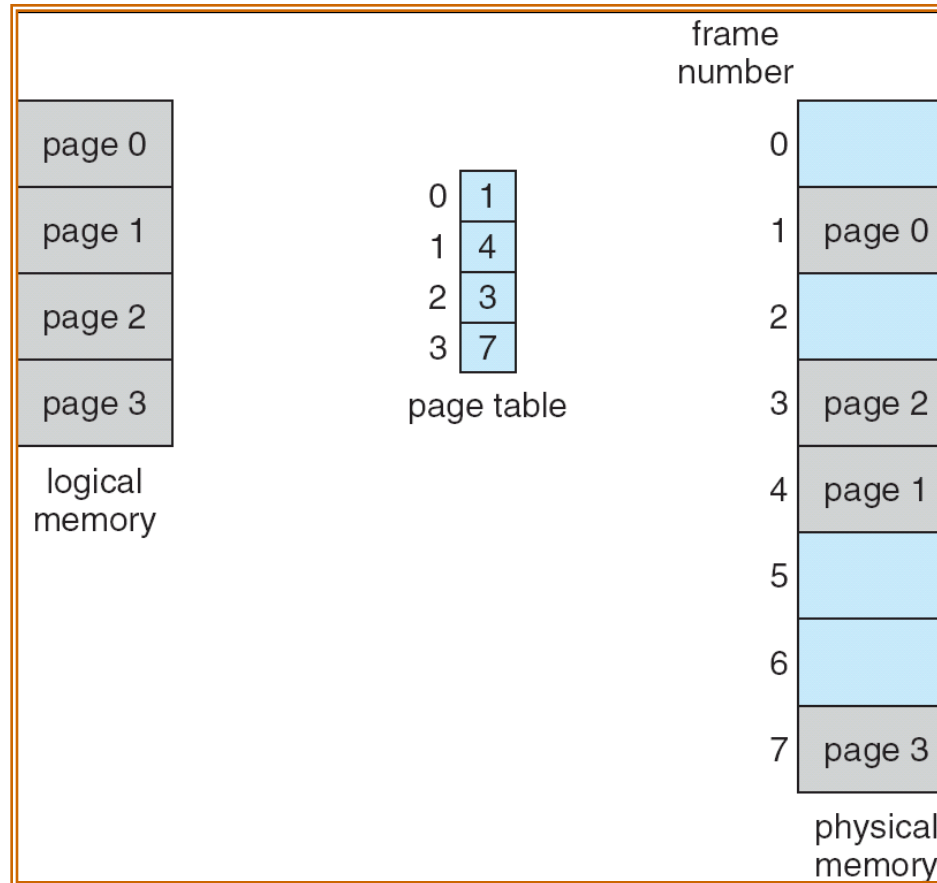
# Paging (2/11)

- The programmer/user sees a process as a contiguous block
  - Think of arrays, adjacent instructions, ...
- What if holes are used totally or partially in physical memory to store a process ?
  - We would need to keep a list of where each of the different logical parts are stored !
  - Too complicated ! Time/Space consuming !

# Paging : Basic Scheme (3/11)

1. Break physical memory into fixed-size blocks called **frames**
2. Break logical memory into fixed-size blocks called **pages**
3. Frames and pages are of identical size
4. Address translation is done using a **page table**

# Paging : Basic Scheme (4/11)

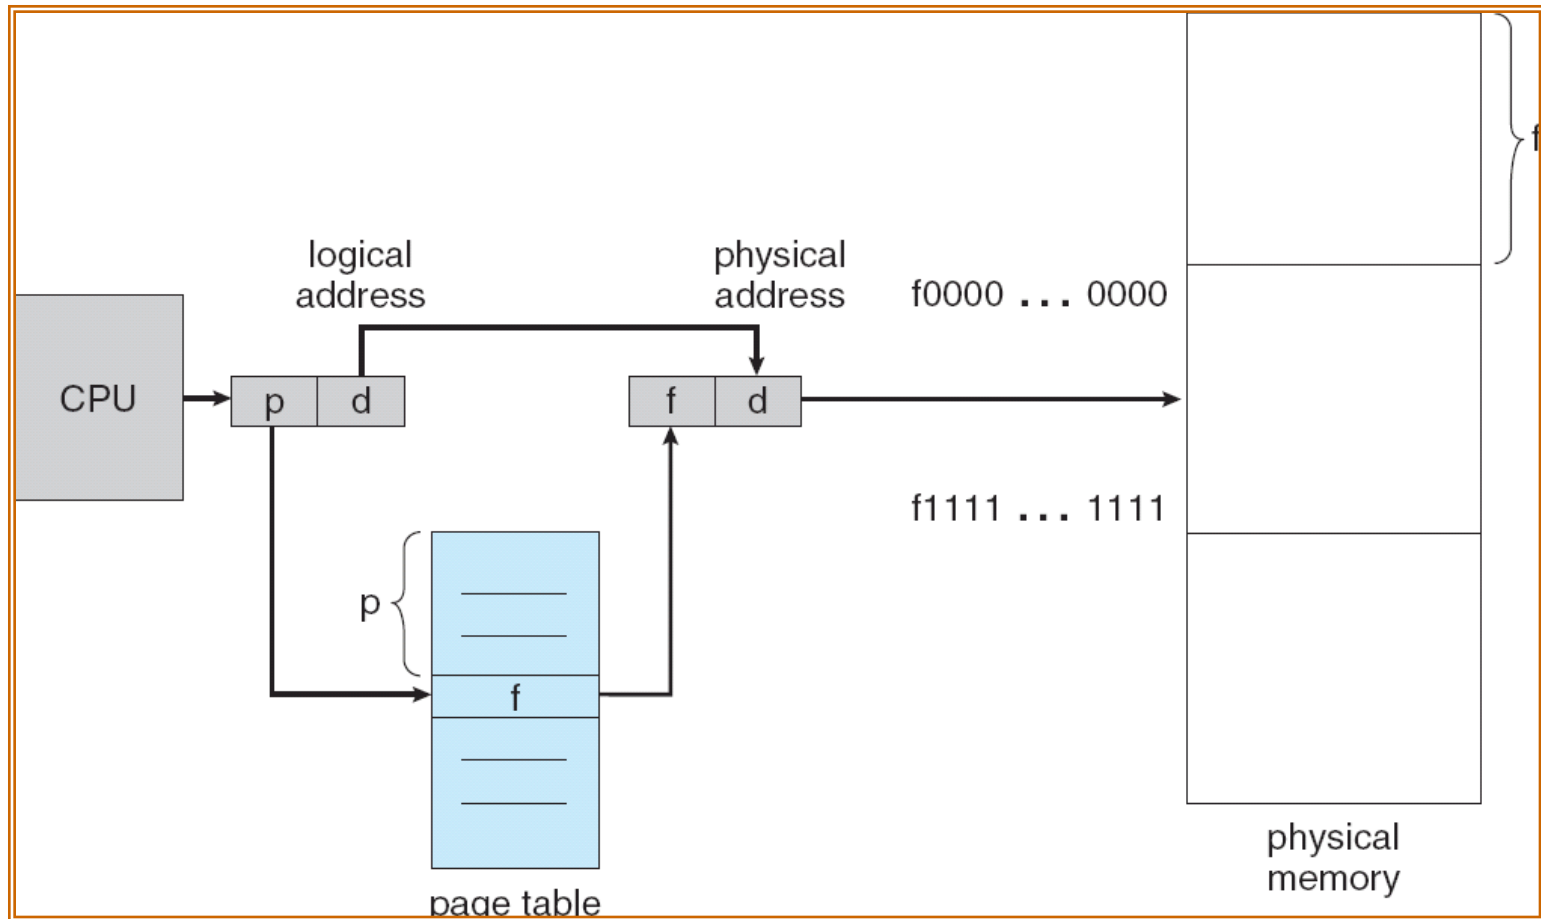


# Paging : Basic Scheme (5/11)

- Address Translation scheme
  - Address generated by the CPU is divided in 2 parts
    - Page number : used as an index on the page table
    - Offset : address inside the page
  - Every valid entry in the page table holds the address of its corresponding physical frame



# Paging : Basic Scheme (6/11)



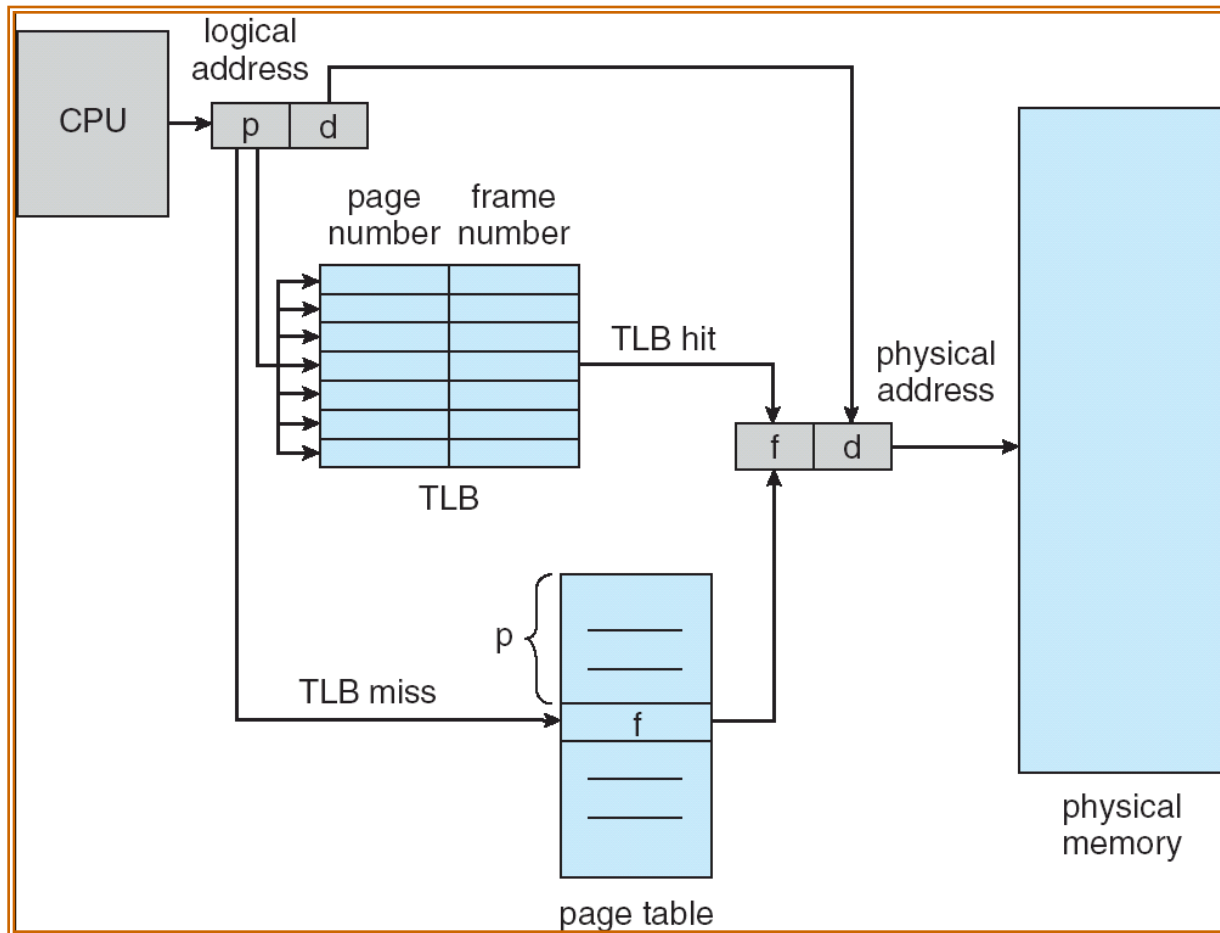
# Paging (7/11)

- Advantages
  - Solves the EXTERNAL fragmentation problem
  - Shared pages
- Drawbacks
  - Internal fragmentation : on average, one half page per process
    - Small page sizes are desirable BUT...what about the page table size ?
  - **One more** memory access to the page table in order to produce the physical address => 2 times slower !
    - Address cache : Translation Lookaside Buffer (TLB)

# Paging (8/11)

- TLB
  - Associative high-speed memory
  - Its benefits are based on the principles of spatial and temporal locality

# Paging with TLB (9/11)



# Hierarchical Paging (10/11)

- Large logical address space => large page table size
  - $2^{64}$  with a page size of 4KB => page table of  $2^{52}$  entries !
- Use an N-level paging algorithm

# Hierarchical Paging (11/11)

- 2-level paging scheme
  - Break the logical address space into Blocks
    - Break blocks into sub-blocks

